

Titanic Classification

Importing the Dependencies

```
In [80]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

```
In [81]: #Load the data from csv file to Pandas DataFrame
df=pd.read_csv('C:/Users/Harshita/OneDrive/Desktop/Titanic-dataset.csv')
```

```
In [82]: # printing the first 5 rows of the dataframe
df.head()
```

Out[82]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	C
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	

```
In [83]: # number of rows and Columns
df.shape
```

Out[83]: (891, 12)

```
In [84]: # getting some informations about the data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   PassengerId     891 non-null    int64
 1   Survived        891 non-null    int64
 2   Pclass          891 non-null    int64
 3   Name            891 non-null    object
 4   Sex             891 non-null    object
 5   Age             714 non-null    float64
 6   SibSp           891 non-null    int64
 7   Parch           891 non-null    int64
 8   Ticket          891 non-null    object
 9   Fare            891 non-null    float64
10   Cabin           204 non-null    object
11   Embarked        889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
In [85]: # check the number of missing values in each column
df.isnull().sum()
```

```
Out[85]: PassengerId     0
Survived               0
Pclass                 0
Name                   0
Sex                    0
Age                   177
SibSp                  0
Parch                  0
Ticket                 0
Fare                   0
Cabin                 687
Embarked               2
dtype: int64
```

Handling the Missing Value

```
In [86]: df = df.drop(columns='Cabin', axis=1)
```

```
In [87]: df['Age'].fillna(df['Age'].mean(), inplace=True)
```

```
In [88]: print(df['Embarked'].mode())
```

```
0    S
Name: Embarked, dtype: object
```

```
In [89]: print(df['Embarked'].mode()[0])
```

```
S
```

```
In [90]: df['Embarked'].fillna(df['Embarked'].mode()[0], inplace=True)
```

```
In [91]: df.isnull().sum()
```

```
Out[91]: PassengerId    0
Survived              0
Pclass               0
Name                 0
Sex                  0
Age                  0
SibSp                0
Parch                0
Ticket              0
Fare                 0
Embarked             0
dtype: int64
```

```
In [92]: # getting some statistical measures about the data
df.describe()
```

```
Out[92]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	891.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	257.353842	0.486592	0.836071	13.002015	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	22.000000	0.000000	0.000000	7.910400
50%	446.000000	0.000000	3.000000	29.699118	0.000000	0.000000	14.454200
75%	668.500000	1.000000	3.000000	35.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200


```
In [93]: df['Survived'].value_counts()
```

```
Out[93]: 0    549
1     342
Name: Survived, dtype: int64
```

Data Visualization


```
In [94]: sns.countplot(x='Survived', data=df)
```

```
Out[94]: <Axes: xlabel='Survived', ylabel='count'>
```



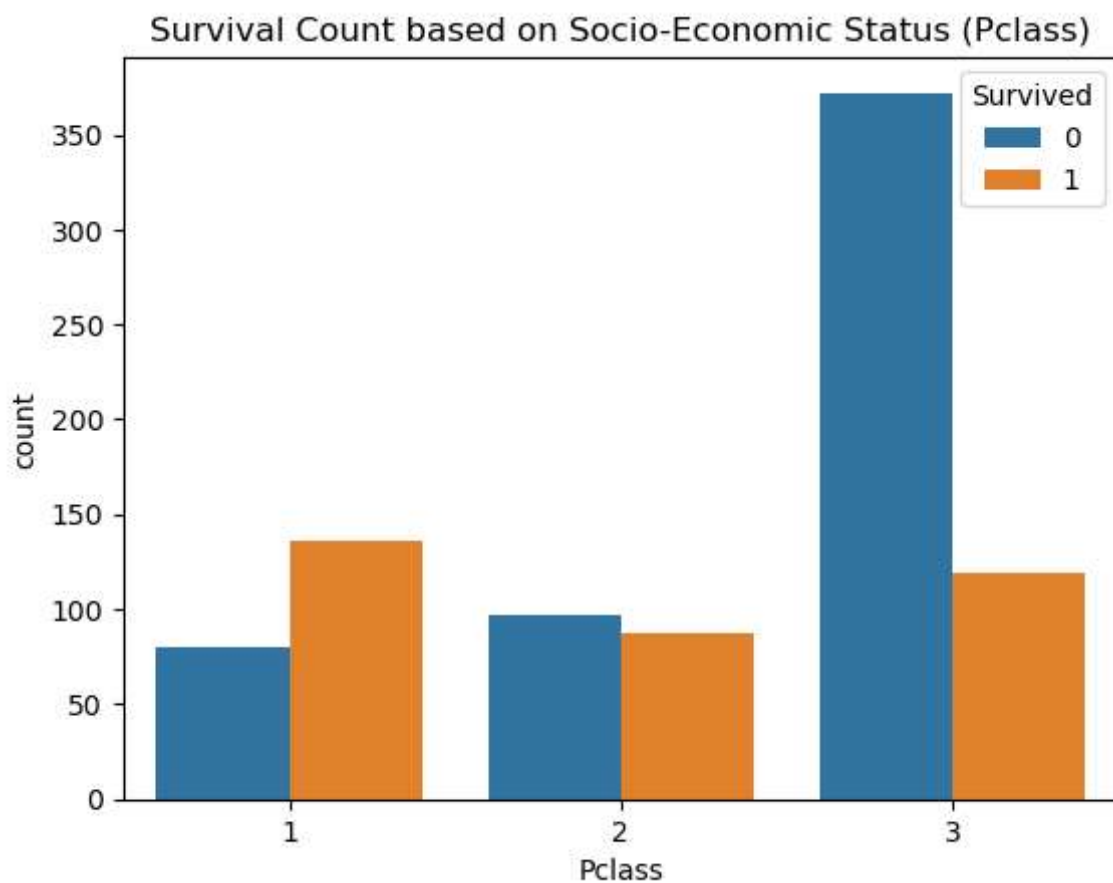
```
In [95]: sns.countplot(x='Sex',hue='Survived',data=df)  
plt.title('Survival Count based on Gender')
```

```
Out[95]: Text(0.5, 1.0, 'Survival Count based on Gender')
```




```
In [96]: sns.countplot(x='Pclass',hue='Survived', data=df)  
plt.title('Survival Count based on Socio-Economic Status (Pclass)')
```

Out[96]: Text(0.5, 1.0, 'Survival Count based on Socio-Economic Status (Pclass)')



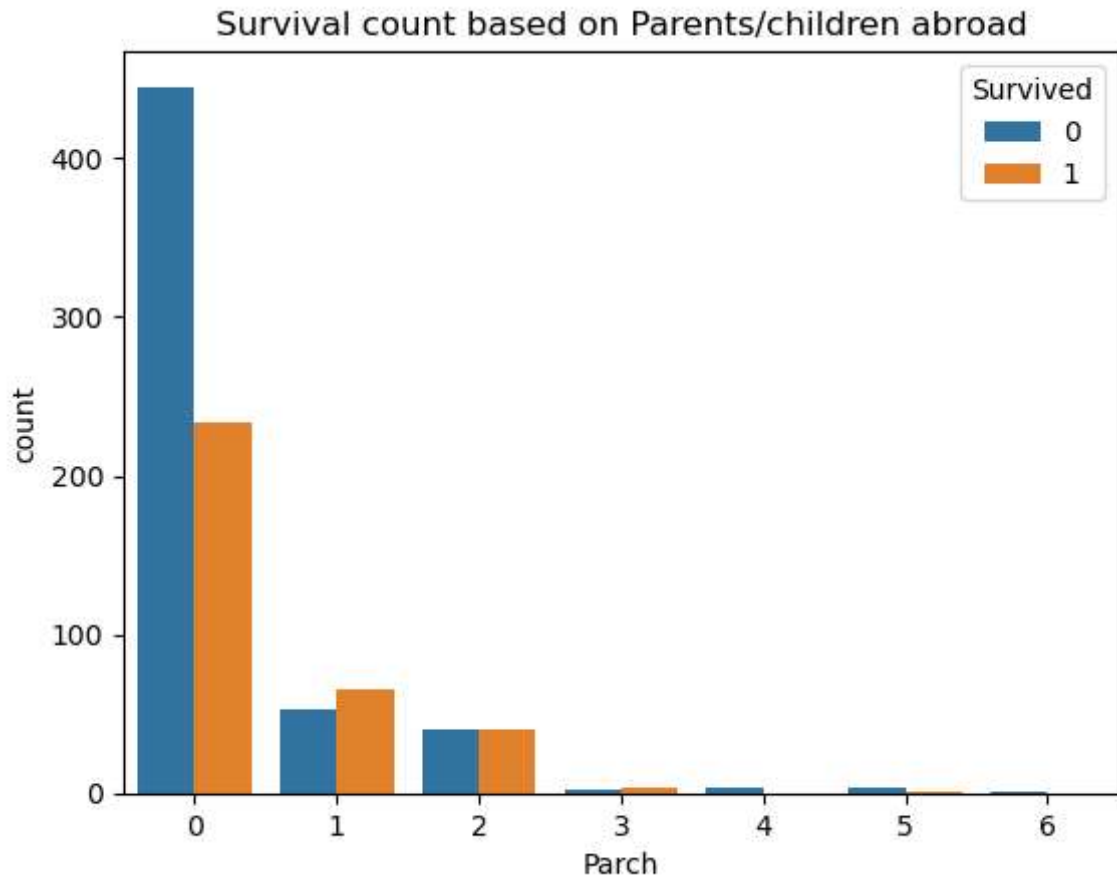
```
In [97]: sns.histplot(x='Age', hue='Survived', data=df, kde=True, palette='coolwarm')  
plt.title('Survival Count based on Age')
```

Out[97]: Text(0.5, 1.0, 'Survival Count based on Age')



```
In [98]: sns.countplot(x='Parch',hue='Survived',data=df)
plt.title('Survival count based on Parents/children abroad')
```

```
Out[98]: Text(0.5, 1.0, 'Survival count based on Parents/children abroad')
```



Encoding the Categorical Columns

```
In [99]: df['Embarked'].value_counts()
```

```
Out[99]: S    646
         C    168
         Q     77
         Name: Embarked, dtype: int64
```

```
In [100]: df['Sex'].value_counts()
```

```
Out[100]: male    577
         female  314
         Name: Sex, dtype: int64
```

```
In [101]: df.replace({'Sex':{'male':0,'female':1}, 'Embarked':{'S':0,'C':1,'Q':2}}, inplace=True)
```



```
In [102]: df.head()
```

```
Out[102]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Emb
0	1	0	3	Braund, Mr. Owen Harris	0	22.0	1	0	A/5 21171	7.2500	
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	1	38.0	1	0	PC 17599	71.2833	
2	3	1	3	Heikkinen, Miss. Laina	1	26.0	0	0	STON/O2. 3101282	7.9250	
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	1	35.0	1	0	113803	53.1000	
4	5	0	3	Allen, Mr. William Henry	0	35.0	0	0	373450	8.0500	

Separating features & Target

```
In [103]: X = df.drop(columns = ['PassengerId', 'Name', 'Ticket', 'Survived'], axis=1)
Y = df['Survived']
```

```
In [104]: print(X)
```

```

      Pclass  Sex      Age  SibSp  Parch      Fare  Embarked
0         3    0  22.000000      1     0    7.2500         0
1         1    1  38.000000      1     0   71.2833         1
2         3    1  26.000000      0     0    7.9250         0
3         1    1  35.000000      1     0   53.1000         0
4         3    0  35.000000      0     0    8.0500         0
..      ...  ...      ...  ...     ...      ...      ...
886        2    0  27.000000      0     0   13.0000         0
887        1    1  19.000000      0     0   30.0000         0
888        3    1  29.699118      1     2   23.4500         0
889        1    0  26.000000      0     0   30.0000         1
890        3    0  32.000000      0     0    7.7500         2

```

```
[891 rows x 7 columns]
```

In [105]: `print(Y)`

```
0      0
1      1
2      1
3      1
4      0
..
886    0
887    1
888    0
889    1
890    0
Name: Survived, Length: 891, dtype: int64
```

Splitting the data into training data & Test data

In [106]: `X_train, X_test, Y_train, Y_test = train_test_split(X,Y, test_size=0.2, random`

In [107]: `print(X.shape, X_train.shape, X_test.shape)`

```
(891, 7) (712, 7) (179, 7)
```

Model training

In [108]: `model = LogisticRegression()`

In [109]: `# training the Logistic Regression model with training data`
`model.fit(X_train, Y_train)`

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\linear_model_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (`max_iter`) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
`n_iter_i = _check_optimize_result(`

Out[109]:

```
LogisticRegression
LogisticRegression()
```

```
In [110]: # accuracy on training data
X_train_prediction = model.predict(X_train)
```

```
In [111]: print(X_train_prediction)
```

```
[0 1 0 0 0 0 0 1 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 0 1 0 0 1 0 1 1 0 0 1 0 1
 0 0 0 0 0 0 1 1 0 0 1 0 1 0 1 0 0 0 0 0 0 1 0 1 0 0 1 1 0 0 1 1 0 1 0 0 1
 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 0 1 0 0 0 1 1 1 0 1 0 0 0 0 0 1 0 0 0
 1 1 0 0 1 0 0 1 0 0 1 0 0 1 0 1 0 1 0 1 0 1 1 1 1 1 1 0 0 1 1 1 0 0 1 0 0
 0 0 0 0 1 0 1 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 1 0 1 0 1 1 1
 0 0 0 1 0 0 0 1 0 0 1 0 0 0 1 1 0 1 0 0 0 0 0 0 1 1 0 1 1 1 1 0 0 0 0 0 0 0
 0 1 0 0 1 1 1 0 0 1 0 1 1 1 0 0 1 0 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 1 0 0 0
 0 0 0 0 0 0 1 0 1 0 0 1 0 0 1 0 1 0 1 1 0 0 0 0 1 0 1 0 0 1 0 0 0 1 0 0 0
 0 1 1 0 0 0 0 0 0 1 0 1 0 0 0 0 0 1 1 1 0 0 0 0 1 0 1 0 0 0 0 0 0 1 1 0 1 1
 0 1 1 1 0 0 0 0 0 0 0 0 0 0 1 0 0 1 1 1 0 1 0 0 0 0 1 1 0 0 0 1 0 1 1 1 0 0
 0 0 1 0 0 0 1 1 0 0 1 0 0 0 0 1 0 0 0 0 0 1 0 0 0 0 1 0 1 1 1 0 1 1 0 0 0
 0 1 0 1 0 0 1 1 0 0 0 0 1 0 0 0 0 1 1 0 1 0 1 0 0 0 0 0 1 0 0 0 0 1 1 0 0
 1 0 1 0 0 1 0 0 0 0 0 0 0 0 1 0 0 1 1 0 0 0 0 1 1 0 1 0 0 1 0 0 0 1 1 0 1 0
 0 0 0 0 1 0 0 1 0 1 1 0 0 1 0 0 1 0 0 0 1 0 1 1 0 0 1 1 0 1 0 1 1 1 0 1 0
 0 1 0 0 1 0 0 1 0 0 0 0 1 1 0 0 1 0 1 0 0 0 0 0 0 1 1 1 0 0 1 1 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 0 0 0 0
 0 0 1 0 0 0 0 0 1 0 1 0 1 0 0 0 1 0 1 1 1 0 0 0 1 0 1 0 0 0 1 1 1 0 0 1 1
 0 0 0 1 0 1 0 0 0 0 0 1 1 0 1 1 1 0 0 0 1 0 0 0 0 1 0 0 0 1 0 0 1 0 0 0 0
 1 0 0 1 0 1 0 0 0 1 1 1 1 1 0 0 1 1 0 1 1 1 1 0 0 0 1 1 0 0 1 0 0 0 0 0 0
 0 0 0 1 1 0 0 1 0]
```

Model Evaluation

Logistic Regression

```
In [112]: training_data_accuracy = accuracy_score(Y_train, X_train_prediction)
print('Accuracy score of training data : ', training_data_accuracy)
```

Accuracy score of training data : 0.8075842696629213

```
In [113]: # accuracy on test data
X_test_prediction = model.predict(X_test)
```

```
In [114]: print(X_test_prediction)
```

```
[0 0 1 0 0 0 0 0 0 0 0 0 1 1 0 0 1 0 0 1 0 1 1 0 1 0 1 1 0 0 0 0 0 0 0 1 1
 0 0 0 0 0 1 0 0 1 1 0 0 1 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 1 0 0 0 1 0 1 0
 1 0 0 0 1 0 1 0 0 0 1 1 0 0 1 0 0 0 0 0 0 1 0 1 0 0 1 0 1 1 0 1 1 0 0 0 0
 0 0 0 1 1 0 1 0 0 1 0 0 0 0 0 0 0 1 0 0 0 0 1 1 0 0 0 0 0 0 1 1 1 1 0 1 0 0
 0 1 0 0 0 0 1 0 0 1 1 0 1 0 0 0 1 1 0 0 1 0 0 1 1 1 0 0 0 0 0]
```

```
In [115]: test_data_accuracy = accuracy_score(Y_test, X_test_prediction)
print('Accuracy score of test data : ', test_data_accuracy)
```

Accuracy score of test data : 0.7821229050279329